

# NoSQL数据库管理

**NSD    NoSQL**

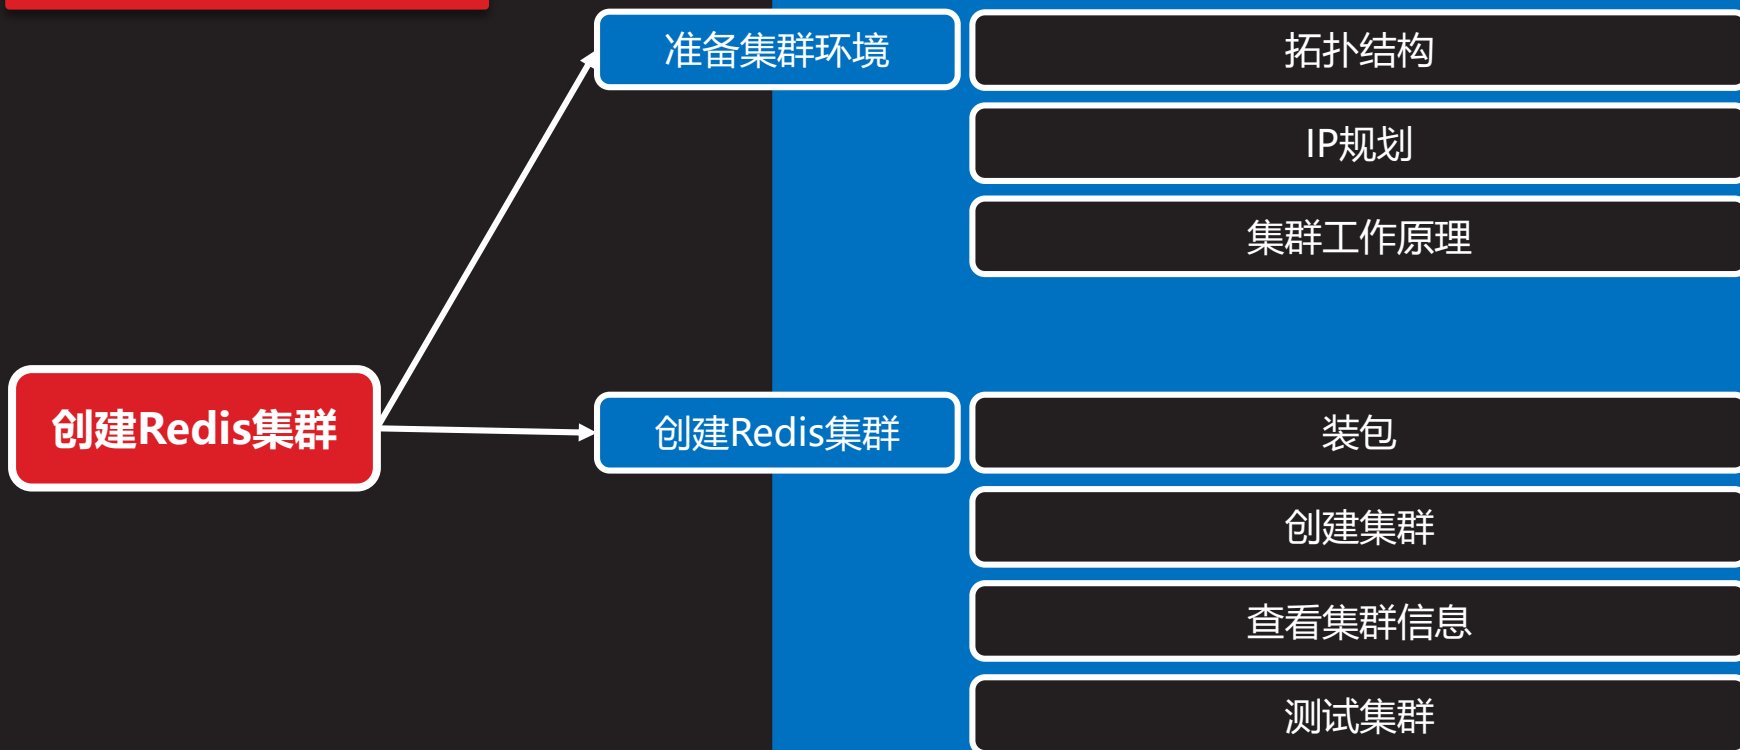
**DAY02**

# 内容

上午	09:00 ~ 09:30	作业讲解和回顾
	09:30 ~ 10:20	创建Redis集群
	10:30 ~ 11:20	
	11:30 ~ 12:00	管理集群
下午	14:00 ~ 14:50	
	15:00 ~ 15:50	
	16:10 ~ 17:00	
	17:10 ~ 18:00	总结和答疑



# 创建Redis集群

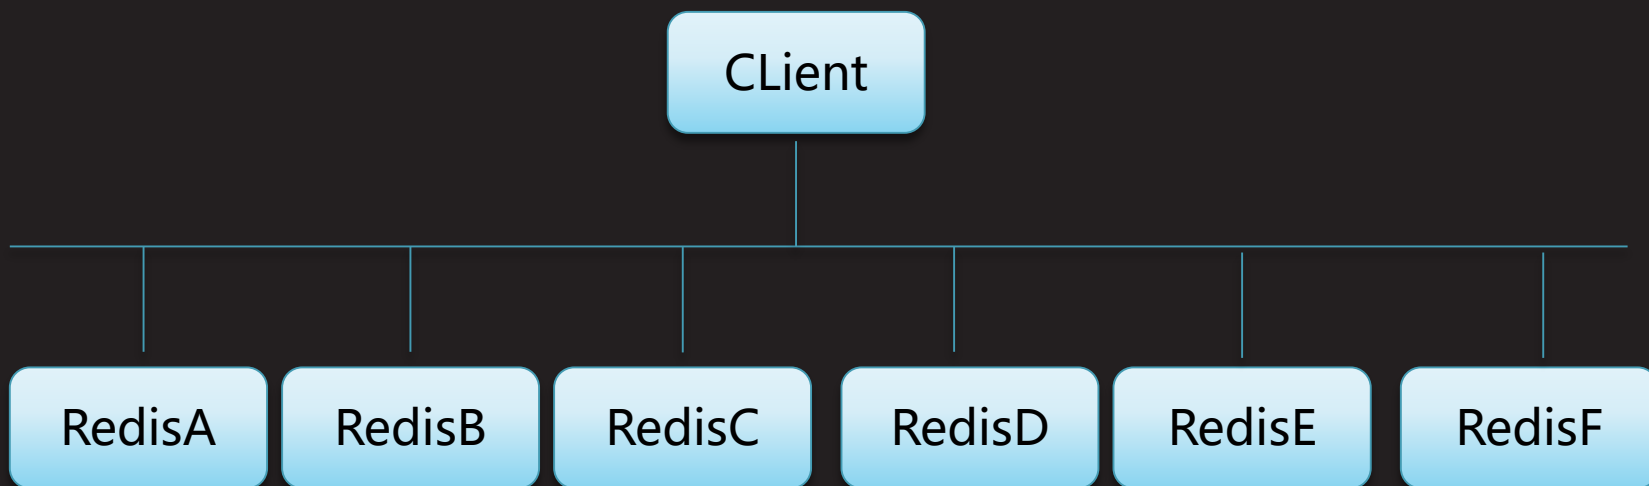


# 准备集群环境



# 拓扑结构

- 六台Redis服务器



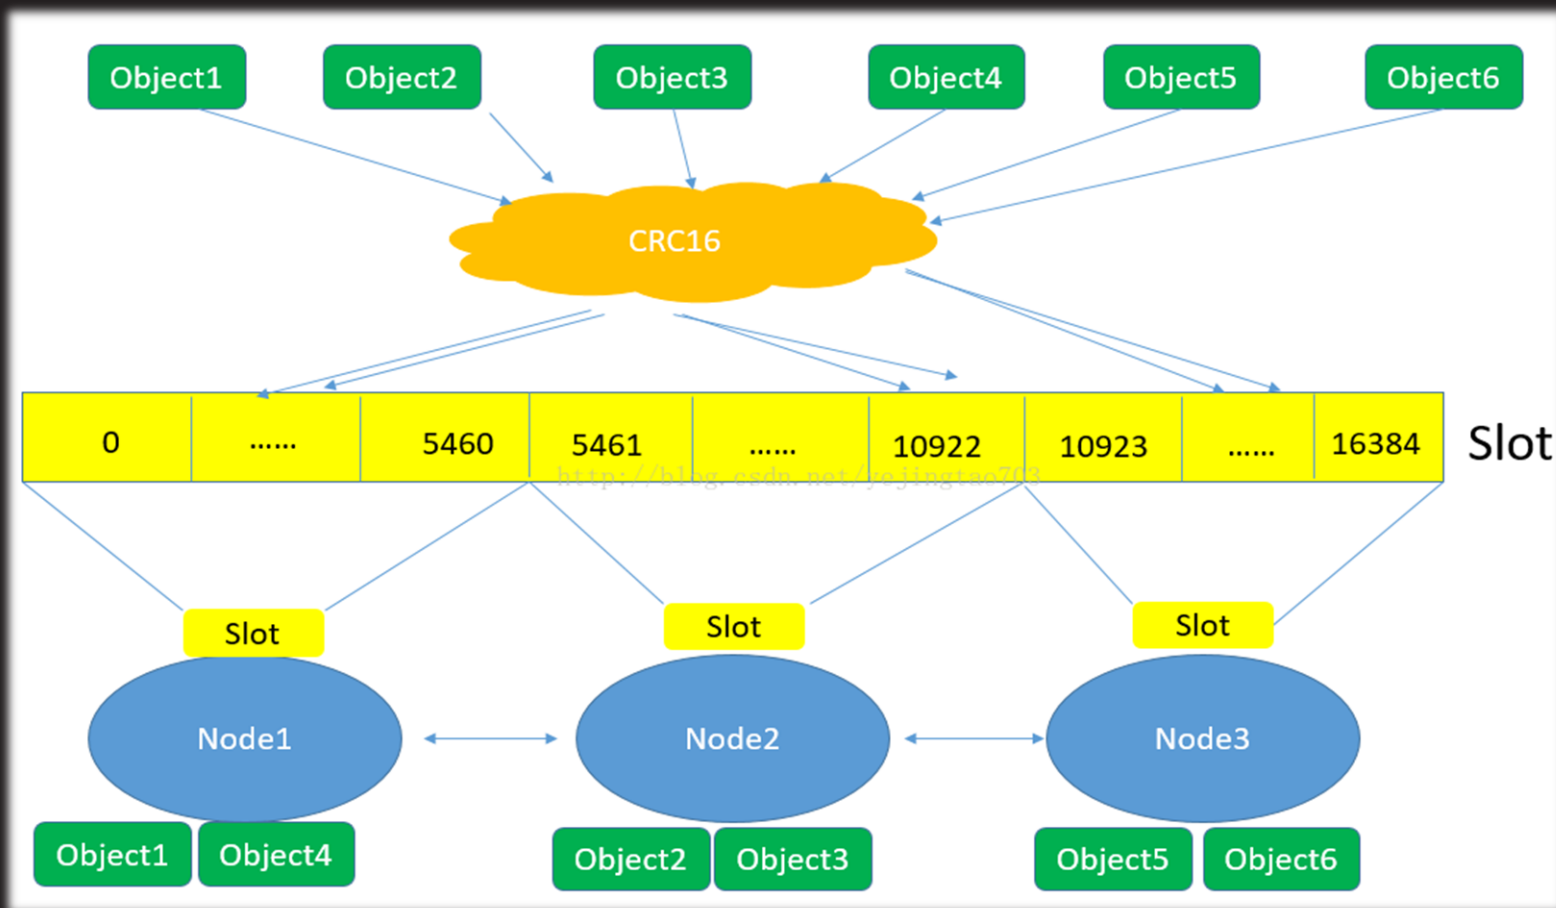
# IP地址规划

- redis 服务器 IP 地址及端口规划
  - redisA 192.168.4.51 6351
  - redisB 192.168.4.52 6352
  - redisC 192.168.4.53 6353
  - redisD 192.168.4.54 6354
  - redisE 192.168.4.55 6355
  - redisF 192.168.4.56 6356



# 集群工作原理

知识讲解



# 创建Redis集群

---



# 装包

- 为6台主机安装并运行redis服务

```
# yum -y install gcc gcc-c++
# tar -zxvf redis-4.0.8.tar.gz
# cd redis-4.0.8/
# make
# make install
# ./utils/install_server.sh
```

- 调整配置文件

```
# vim /etc/redis/redis.conf
bind IP地址
daemonize yes
port xxxx
cluster-enabled yes
cluster-config-file nodes-xxxx.conf
cluster-node-timeout 5000
```

```
//只写物理接口IP地址
//守护进程方式运行
//端口号不要使用默认的6379
//启用集群
//指定集群信息文件
//请求超时 5 秒
```



# 装包 (续1)

- 查看服务信息

```
[root@host51 ~]# /etc/init.d/redis_6379 status
```

```
Redis is running (21201)
```

```
[root@host51 ~]#
```

```
[root@host51 ~]# netstat -utnlp | grep redis-server
```

```
tcp      0      0 192.168.4.51:6351    0.0.0.0:*        LISTEN
21201/redis-server
```

```
tcp      0      0 192.168.4.51:16351   0.0.0.0:*        LISTEN
21201/redis-server
```

```
[root@host51 ~]#
```

集群通信端口 默认服务端口+10000



## 装包 (续2)

- 查看集群信息

```
[root@host51 ~]# redis-cli -h 192.168.4.51 -p 6351
192.168.4.50:6350> cluster info
cluster_state:fail
cluster_slots_assigned:0
cluster_slots_ok:0
cluster_slots_pfail:0
cluster_slots_fail:0
cluster_known_nodes:1
cluster_size:0
.....
192.168.4.50:6350>
192.168.4.50:6350> cluster nodes
e081313ec843655d9bc5a17f3bed3de1dcc1d2b
192.168.4.50:6350@16350 myself, master - 0 1530781129976 11
connected
192.168.4.50:6350>
```



# 创建集群

- 在选中的一台redis服务器上，执行创建集群脚本
  - 部署ruby脚本运行环境
  - 创建集群

```
# yum -y install ruby rubygems
# rpm -ivh --nodeps ruby-devel-2.0.0.648-30.el7.x86_64.rpm
# gem install redis-3.2.1.gem

# cd redis-3.2.0/src/
# ./redis-trib.rb create --replicas 1 host:port host:port .....
```

--replicas 1 ， 自动为每一个master节点分配一个slave节点



# 创建集群 (续1)

- 创建集群

```
[root@host51 ~]# redis-trib.rb create --replicas 1 \
192.168.4.51:6351 192.168.4.52:6352 \
192.168.4.53:6353 192.168.4.54:6354 \
192.168.4.55:6355 192.168.4.56:6356 \
>>> Creating cluster
>>> Performing hash slots allocation on 6 nodes...
Using 3 masters:
192.168.4.51:6351
192.168.4.52:6352
192.168.4.53:6353
.....
[OK] All nodes agree about slots configuration.
>>> Check for open slots...
>>> Check slots coverage...
[OK] All 16384 slots covered.
```



# 查看集群信息

- 任意一台主机访问本机的redis服务查看即可

> cluster info //查看集群信息

> cluster nodes //查看集群节点信息

```
# redis-cli -h 192.168.4.52 -p 6352
```

```
192.168.4.52:6352> cluster info
```

```
cluster_state:ok
```

```
cluster_slots_assigned:16384
```

```
cluster_slots_ok:16384
```

```
...
```

```
cluster_known_nodes:6
```

```
cluster_size:3
```

```
...
```



# 查看集群信息（续1）

- 任意一台主机访问本机的redis服务查看即可

> cluster info //查看集群信息

> cluster nodes //查看集群节点信息

```
# redis-cli -h 192.168.4.52 -p 6352
192.168.4.52:6352> cluster nodes
6e841e2610c3d4d..... 192.168.4.51:6351@16351 master .....
2de5136be52a327..... 192.168.4.54:6354@16354 slave .....
3198014263d26a6..... 192.168.4.56:6356@16356 master .....
82781de818fb83cc..... 192.168.4.53:6353@16353 master .....
b52feba0e6422b44..... 192.168.4.52:6352@16352 myself,slave .....
a0c8e18619828487..... 192.168.4.55:6355@16355 slave .....
... ..
```



# 测试集群

- 在客户端访问任意一台master主机存数据

– redis-cli -c -h ip地址 -p 端口

> set key values //存数据

```
[root@host51 ~]# redis-cli -c -h 192.168.4.51 -p 6351
192.168.4.51:6351> set school tarena
-> Redirected to slot [8455] located at 192.168.4.52:6352
OK
192.168.4.52:6352> set class linux
OK
192.168.4.52:6352> set pay 26800
-> Redirected to slot [4013] located at 192.168.4.51:6351
OK
```





# 测试集群 (续1)

- 在客户端访问任意一台master主机取数据

– redis-cli -c -h ip地址 -p 端口

> get key //取数据

```
[root@host51 ~]# redis-cli -c -h 192.168.4.53 -p 6353
192.168.4.53:6353> get name
-> Redirected to slot [5798] located at 192.168.4.52:6352
"bob"
192.168.4.52:6352> keys *
1) "school"
2) "name"
3) "class"
4) "age"
192.168.4.52:6352> get age
"19"
```



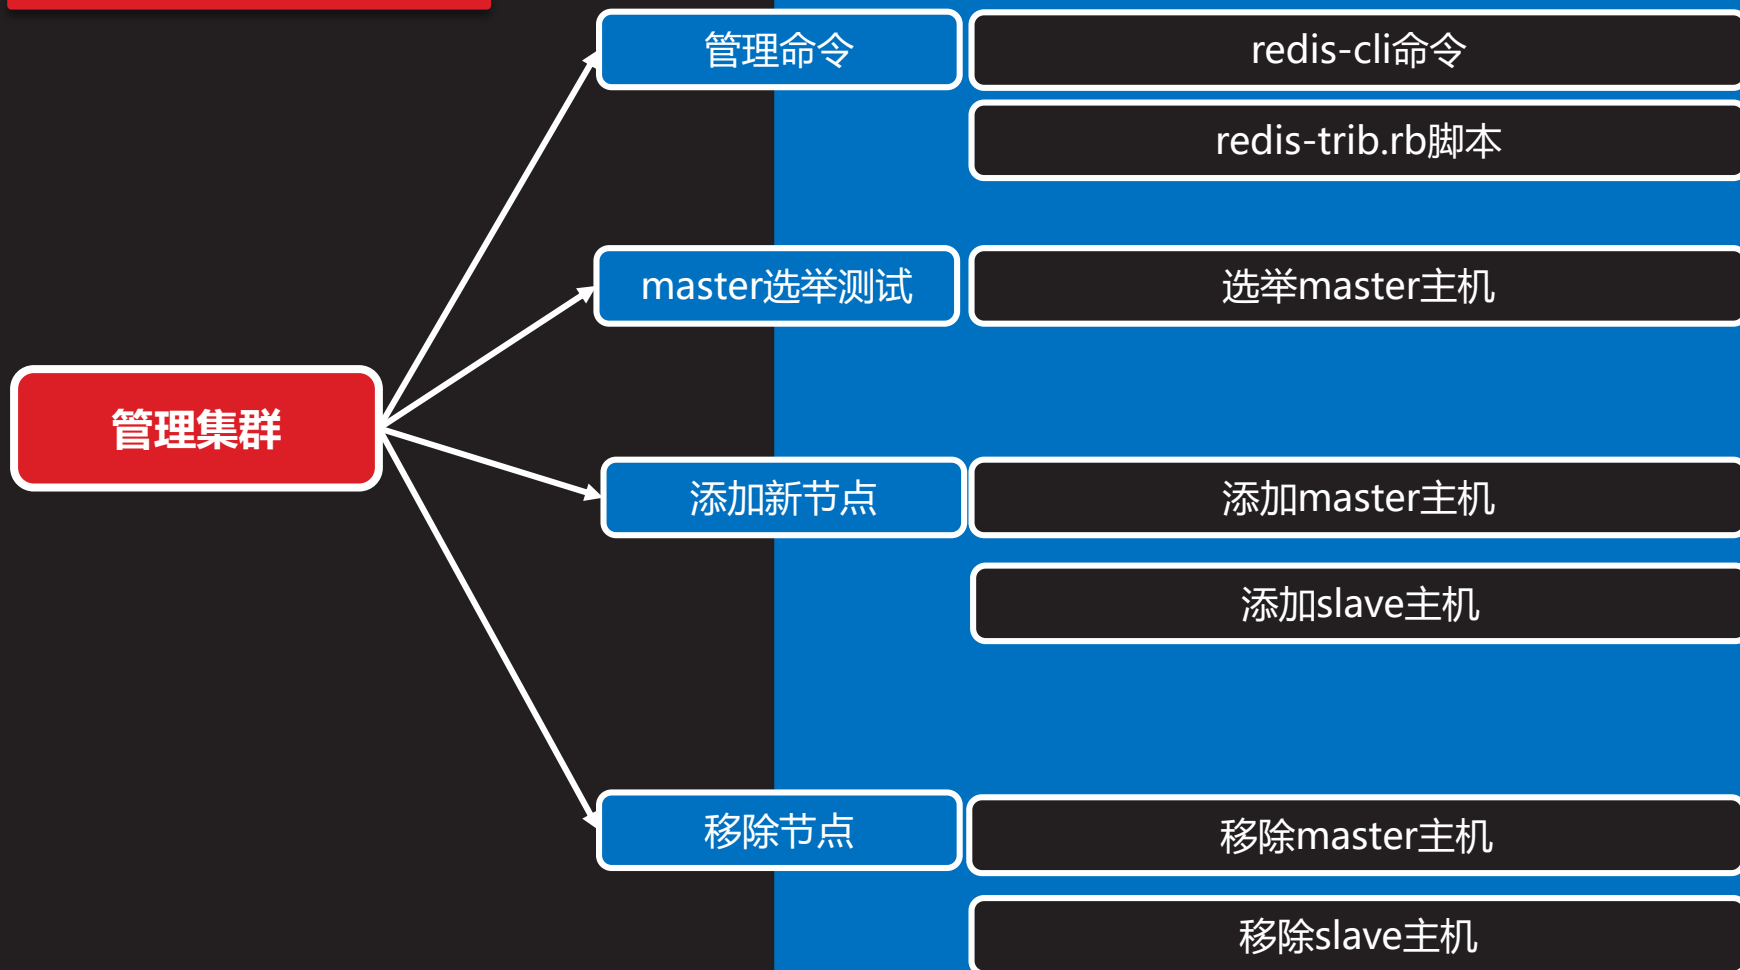
# 案例1：部署redis集群

具体要求如下：

- 准备集群环境
- 安装redis并创建集群
- 查看集群信息



# 管理集群



# 管理命令

---

# redis-cli命令

- 查看命令帮助
  - redis-cli -h
- 常用选项
  - -h IP地址
  - -p 端口
  - -c 集群模式



# redis-trib.rb脚本

- 语法格式
  - redis-trib.rb 选项 参数
- 选项
  - add-node 添加master主机
  - check 检测集群
  - reshard 重新分片
  - add-node --slave 添加slave主机
  - del-node 删除主机



# master选举测试

---

# 选举master主机

- 停止master 主机的 Redis服务
  - master宕机后对应的slave自动被选举为master
  - 原master启动后 会自动配置为当前master的slave
- 查看集群主机信息
  - `redis-cli -h master_ip -p master_port`

```
[root@host51 ~]# redis-cli -h 192.168.4.51 -p 6351  
192.168.4.51:6351> cluster info
```





# 添加新节点

---

# 添加master主机

- 部署一台新redis服务器
  - 装包
  - 初始化
  - 启用集群配置
  - 重启服务
- 添加master主机步骤
  - 添加master主机
  - 检查主机
  - 重新分片



# 添加master主机 (续1)

- 添加master主机
  - 添加时不指定主机角色，默认新主机被选为master
  - # ./redis-trib.rb add-node 新主机Ip:端口

192.168.4.51:6351

```
[root@host51 ~]# redis-trib.rb add-node 192.168.4.50:6350
192.168.4.51:6351
```

```
>>> Adding node 192.168.4.50:6350 to cluster
192.168.4.51:6351
```

.....

```
[OK] All nodes agree about slots configuration.
```

```
>>> Check for open slots...
```

```
>>> Check slots coverage...
```

```
[OK] All 16384 slots covered.
```

```
>>> Send CLUSTER MEET to node 192.168.4.50:6350 to
make it join the cluster.
```

```
[OK] New node added correctly.
```



## 添加master主机 (续2)

- 检测集群主机
  - ./redis-trib.rb check 192.168.4.51:6351
  - 主机角色为master
  - 无槽位数量

```
[root@host51 ~]# redis-trib.rb check 192.168.4.51:6351
>>> Performing Cluster Check (using node 192.168.4.51:6351)
.....
.....
M: e081313ec843655d9bc5a17f3bed3de1dcc1d2b
192.168.4.50:6350
  slots: (0 slots) master
  0 additional replica(s)
.....
.....
```



## 添加master主机 (续3)

- 重新分片

- ./redis-trib.rb reshard 192.168.4.51:6351
- 指定移出hash槽个数
- 指定接收hash槽主机ID
- 指定移出hash槽主机ID

```
[root@host51 ~]# redis-trib.rb reshard 192.168.4.51:6351
How many slots do you want to move (from 1 to 16384)? 4096
What is the receiving node ID?
```

```
e081313ec843655d9bc5a17f3bed3de1dcc1d2b
```

```
Please enter all the source node IDs.
```

```
Type 'all' to use all the nodes as source nodes for the hash slots.
```

```
Type 'done' once you entered all the source nodes IDs.
```

```
Source node #1:3550af084c94889ae2d01103e6da5793fdd851fc
```

```
Source node #2:done
```

```
.....
```

```
Do you want to proceed with the proposed reshard plan (yes/no)? yes
```

# 添加slave主机

- 部署一台新redis服务器

- 装包
- 初始化
- 启用集群配置
- 运行服务

- 添加slave主机

```
# ./redis-trib.rb add-node --slave [ --master-id id值 ]  
ip地址:端口 192.168.4.51:6351
```

如果不指定主节点的id的话，会把新节点 随机添加为 从节点 最少的主的从

# 添加slave主机 (续1)

- 添加slave主机

```
[root@host51 ~]# redis-trib.rb add-node --slave
192.168.4.57:6357 192.168.4.51:6351
>>> Adding node 192.168.4.57:6357 to cluster
192.168.4.51:6351
>>> Performing Cluster Check (using node 192.168.4.51:6351)
.....
.....
[OK] All 16384 slots covered.
Automatically selected master 192.168.4.51:6351
>>> Send CLUSTER MEET to node 192.168.4.57:6357 to make
it join the cluster.
Waiting for the cluster to join.
>>> Configure node as replica of 192.168.4.51:6351.
[OK] New node added correctly.
```



# 移除节点

---



# 移除master主机

- 配置步骤
  - 重新分片释放占用的hash槽
  - 移除master主机

```
# redis-trib.rb reshard 192.168.4.51:6351
```

```
# redis-trib.rb del-node 192.168.4.51:6351 maste主机id值
```



# 移除master主机 (续1)

- 重新分片释放占用的hash槽
  - 指定移出hash槽个数
  - 指定接收hash槽主机ID
  - 指定移出hash槽主机ID

```
[root@host51 ~]# redis-trib.rb reshard 192.168.4.51:6351
How many slots do you want to move (from 1 to 16384)? 4096
What is the receiving node ID?
```

```
4361720c3978aa02347076218580a103c60a6d7f
```

```
Please enter all the source node IDs.
```

```
Type 'all' to use all the nodes as source nodes for the hash slots.
```

```
Type 'done' once you entered all the source nodes IDs.
```

```
Source node #1:e081313ec843655d9bc5a17f3bed3de1dcc1d2b
```

```
Source node #2:done
```

```
.....
```

```
Do you want to proceed with the proposed reshard plan (yes/no)? yes
```



## 移除master主机 (续2)

- 移除master主机

# redis-trib.rb del-node 192.168.4.51:6351 maste主机id值

```
[root@host51 ~]# redis-trib.rb del-node 192.168.4.51:6351 \
e081313ec843655d9bc5a17f3bed3de1dcc1d2b
>>> Removing node
e081313ec843655d9bc5a17f3bed3de1dcc1d2b from cluster
192.168.4.51:6351
>>> Sending CLUSTER FORGET messages to the cluster...
>>> SHUTDOWN the node.
[root@host51 ~]#
```



# 移除slave主机

- 移除slave主机

- 从节点主机没有槽位范围，直接移除即可
- `redis-trib.rb del-node 192.168.4.51:6351 主机id值`

```
# redis-trib.rb del-node 192.168.4.51:6351 \  
f6649ea99b2f01faca26217691222c17a3854381  
>>> Removing node f6649ea99b2f01faca26217691222c17a3854381  
from cluster 192.168.4.51:6351  
>>> Sending CLUSTER FORGET messages to the cluster..  
>>> SHUTDOWN the node.
```



## 案例2：管理redis集群

具体要求如下：

- 练习添加主机
- 练习删除主机



# 总结和答疑

---

管理集群

redis-trib.rb脚本选项总结

总结和答疑

```
graph LR; A[总结和答疑] --> B[管理集群]; B --> C[redis-trib.rb脚本选项总结];
```

# 管理集群



# redis-trib.rb脚本选项总结

- redis-trib.rb 常用选项

选项	作用
create	创建集群
check	检查集群
reshard	重新分片
del-node	删除主机
add-node --slave	添加slave主机
add-node	添加master主机

